

Deep Neural Networks, They Just (Don't) Work

Kushal Jhunjunwala (kushalj@cs)

Andrew Wei (nowei@cs)


Leiyi Zhang (leiyiz@cs)

CSE 499/599G1: Intro to Deep Learning

Autumn 2019

Goal

We wanted to show that DNNs weren't robust and would confidently misclassify when adding targeted noise that would be imperceptible to humans.



What we did

Fast Gradient Sign Method (FGSM)

High level: Take the sign of the gradient and move in the opposite direction



Samoyed (89.36%) classified as Poncho (33.40)%

Projected Gradient Descent (PGD)

High level: Move in the direction opposite the gradient (or towards it if targeted) and project back down to at most an epsilon difference
(See cat picture)

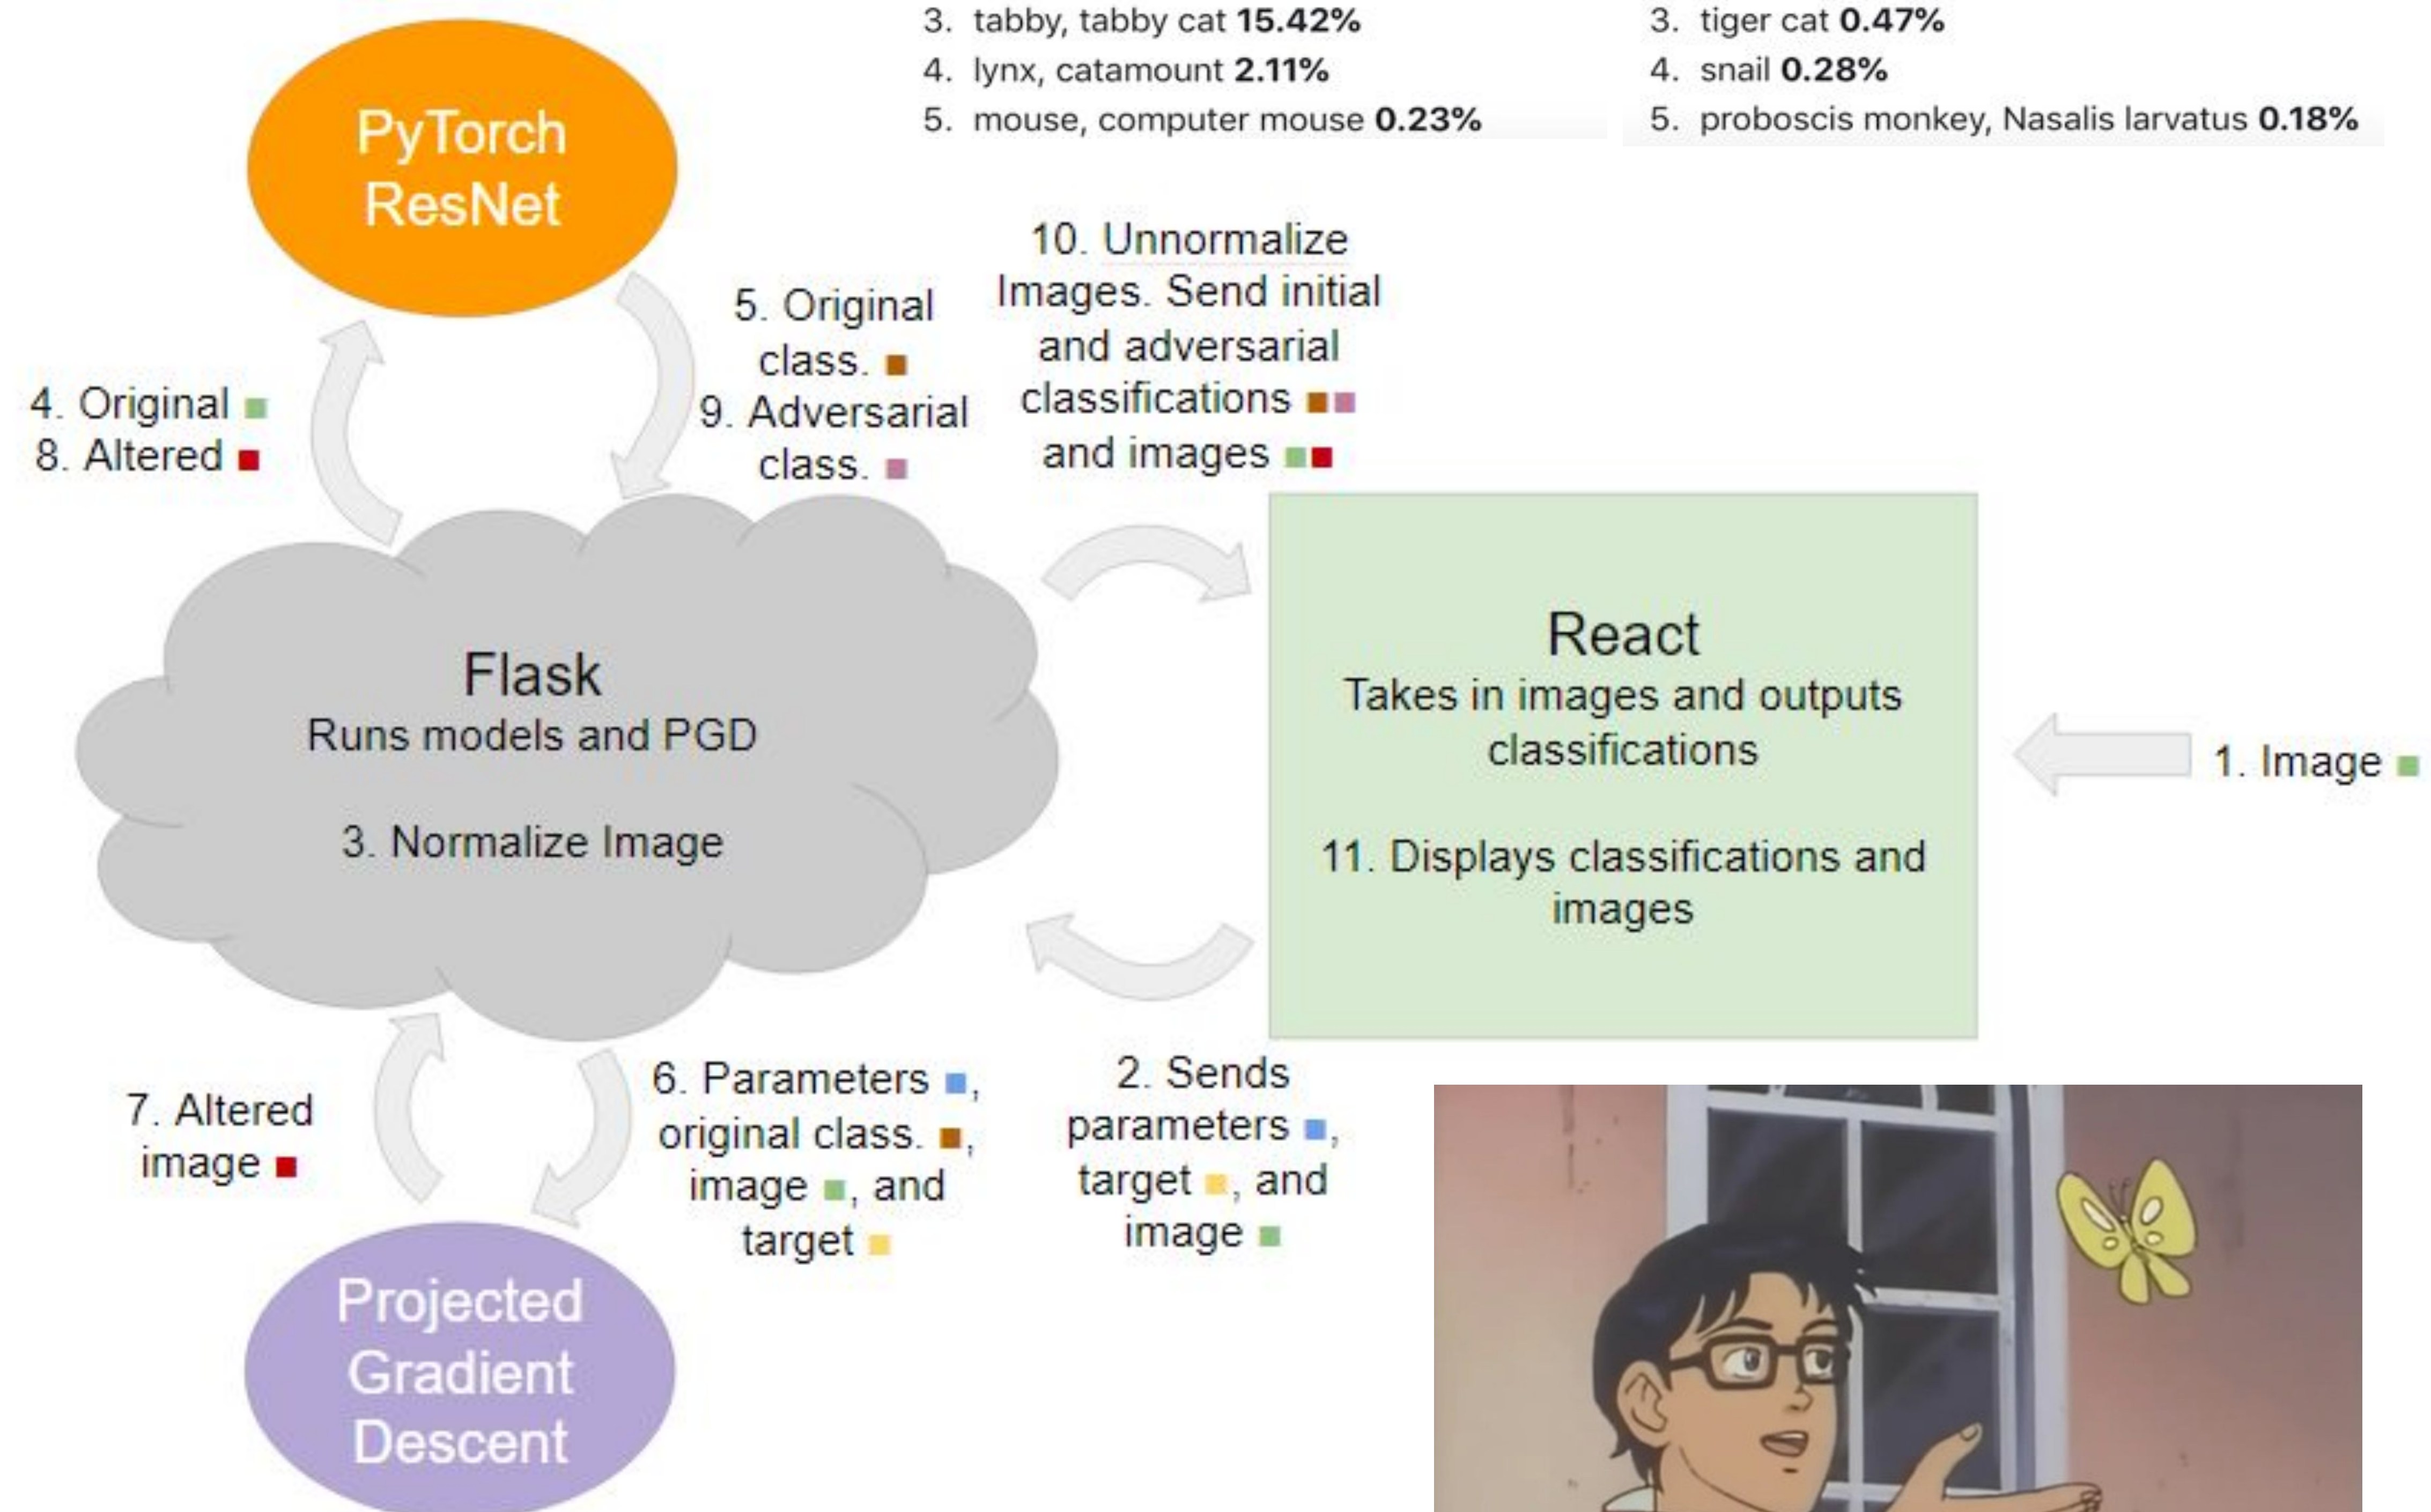
When you ask your NN
why it classified
your cat as a dog:



1. Egyptian cat **41.07%**
2. tiger cat **39.43%**
3. tabby, tabby cat **15.42%**
4. lynx, catamount **2.11%**
5. mouse, computer mouse **0.23%**



1. goldfish, Carassius auratus **94.53%**
2. conch **0.53%**
3. tiger cat **0.47%**
4. snail **0.28%**
5. proboscis monkey, Nasalis larvatus **0.18%**



Problems we faced

- Getting ImageNet data - no longer available through PyTorch
- FGSM perturbed it too much
- Training a robust model
- Perturbations between pixel values get truncated when saving the image

References

- Explaining and Harnessing Adversarial Examples (2014), Ian J. Goodfellow and Jonathon Shlens and Christian Szegedy
- Towards Deep Learning Models Resistant to Adversarial Attacks (2017), Aleksander Madry and Aleksandar Makelov and Ludwig Schmidt and Dimitris Tsipras and Adrian Vladu